

M2 Android Build — Field Test Findings

Observed verdict discrepancies and control cases on the current Android test build

M2 STATUS: OPEN

Structures are encouraging; insect routing is not stable enough to close M2.

This report records observed output versus the behavior already expected in M2. It requests diagnosis, a minimal correction and regression evidence — not an M3 redesign.

POSITIVE CONTROL

Structures / nests

Initial judgments are generally coherent and should be protected from regression.

M2 BLOCKER

Insect verdict routing

Clear or known specimens still reach inconsistent verdict families.

CRITICAL • M2 BLOCKER

Photo 1 — Confirmed Asian hornets, distant group on jar



GROUND TRUTH

Confirmed Asian hornets

CURRENT APP OUTPUT

ORANGE_PROBABLE_NON_CIBLE — 65%

EXPECTED OUTPUT FOR THIS IMAGE

ORANGE_INSUFFISANCE

FINDING

The insects occupy too few pixels for reliable Q1/Q2/Q3 assessment. A cautious insufficiency verdict is justified; a non-target verdict is not. On this confirmed target reference, VERT and NON_CIBLE are prohibited outcomes.

POINT TO INVESTIGATE

Check whether low specimen resolution or the patterned jar lid generated false crabro evidence, then trace the raw observation, normalized traits and final Judge route.

Verdict routing — hornets and clear wasp

Two discrepancies and one useful control pass



PHOTO 2

DISCREPANCY

European hornet under glass

CURRENT APP OUTPUT
ORANGE_INSUFFISANCE — 60%

EXPECTED M2 OUTPUT
ORANGE_PROBABLE_NON_CIBLE

FINDING

Despite reflections, the specimen remains usable enough to identify strong European-hornet / non-target evidence. The existing non-target path did not fire.

POINT TO INVESTIGATE

Compare Gemini's raw observation, the normalized data and the Judge decision. Determine whether the loss occurs upstream or during verdict routing.



PHOTO 3

CONTROL PASS

Same European specimen — successful control

CURRENT APP OUTPUT
ORANGE_PROBABLE_NON_CIBLE — 70%

EXPECTED M2 OUTPUT
ORANGE_PROBABLE_NON_CIBLE

FINDING

The verdict family is correct. Exact species naming is not required here; the important result is a clear non-target warning rather than insufficiency.

CONTROL VALUE

This pass proves the current M2 path can reach the expected outcome. Use it as a comparison point for Photo 2 and explain the variability.




PHOTO 4

DISCREPANCY

Clear wasp-like non-target photo on red background

CURRENT APP OUTPUT
ORANGE_INSUFFISANCE — 60%

EXPECTED M2 OUTPUT
ORANGE_PROBABLE_NON_CIBLE

FINDING

The dorsal image is sharp. Repeated pale transverse abdominal bands and the absence of a single broad orange fourth-tergite signature support a non-target verdict.

POINT TO INVESTIGATE

Verify why the existing wasp / Polistes non-target route requested another image despite clear exclusion evidence in the submitted view.

Hairy-body references — existing VERT pathway

The request is to exercise an M2 rule already present, not to add a new species classifier

EXISTING M2 ROUTE UNDER TEST

Dense, clearly visible body hair → `morphologie_valeur_compacte` / `INSECT_HAIRY_BODY_INCOMPATIBLE` → VERT.




PHOTO 5

HIGH PRIORITY

Scoliid / mammoth wasp — dense hairy body

CURRENT APP OUTPUT

ORANGE_INSUFFISANCE — 60%

EXPECTED M2 OUTPUT

VERT

FINDING

Dense body hair is clearly visible. The paired yellow abdominal spots are additional visible support, but no new Scoliidae-specific rule is requested.

POINT TO INVESTIGATE

Trace why the existing hairy-body incompatibility marker was not emitted or was not honored by the Judge. Correct the smallest failing layer.




PHOTO 6

HIGH PRIORITY

Scoliid / mammoth wasp on flower — strong hair evidence

CURRENT APP OUTPUT

ORANGE_INSUFFISANCE — 60%

EXPECTED M2 OUTPUT

VERT

FINDING

Prominent bristly hair is visible in close-up, with isolated yellow abdominal patches. This should exercise the same existing VERT pathway as Photo 5.

POINT TO INVESTIGATE

Confirm whether the failure comes from visual extraction, normalization or the Judge. Add this image to the M2 regression set after correction.

M2 closure gate

Expected verdict matrix, requested diagnostic evidence and acceptance conditions

Expected verdict matrix

CASE	CURRENT	EXPECTED M2
Photo 1	NON_CIBLE 65%	ORANGE_INSUFFISANCE
Photo 2	INSUFFISANCE 60%	ORANGE_PROBABLE_NON_CIBLE
Photo 3	NON_CIBLE 70%	CONTROL_PASS
Photo 4	INSUFFISANCE 60%	ORANGE_PROBABLE_NON_CIBLE
Photos 5-6	INSUFFISANCE 60%	VERT

Requested diagnosis and delivery

- 1

Reproduce every attached case
Run the same files on the current Android build and record the exact output.
- 2

Provide the decision trace
Raw Gemini response, normalized observation, final Judge verdict and reason code.
- 3

Locate the failing layer
State whether each discrepancy is upstream extraction, normalization or Judge logic.
- 4

Apply the minimal M2 correction
Use existing rules where possible; avoid introducing unrelated architecture changes.
- 5

Deliver a corrected Android build
Retest the six files and include all failures in the M2 regression set.
- 6

Protect the positive branch
Confirm that the correction causes no regression on structures and nests.

ACCEPTANCE GATE

M2 remains open until the corrected build is demonstrated against these cases.

Closure requires the expected verdict family for Photos 1, 2, 4, 5 and 6; continued success on Photo 3; an auditable reason for each result; and no regression on the structure branch.

OUT OF SCOPE FOR THIS REPORT

- M3 proxy / response_schema implementation
- Per-trait experiment or prompt architecture redesign
- A new large blind benchmark or new paid calibration phase